


The Monk Who *Oversold* the AI

*A product leader sold the board on
borrowed intelligence, scaled fast, won
two years of trust, and quietly dissolved
the only moat the product ever had.*

TO Boards, CEOs, heads of product and R&D
FROM Anuj Sadani • Principal SDE, builder of AI systems and the teams behind them
RE Why borrowed AI scales fast and sinks faster



Frontispiece. Borrowed capability scales fast and feels like genius. Then the rented edge converges to zero, and the moat you stopped digging is gone.

Read the **Overture** in five minutes for the whole arc. Read the rest for what the moat actually was, why even a strong ML org gives it away, and how to borrow without dissolving.

CONTENTS

—	Overture: The Two-Year Miracle	<i>the whole arc in five minutes</i>
01	The Company That Had a Moat	<i>R&D as the asset</i>
02	The Reframe	<i>“we don’t need R&D”</i>
03	Why It Worked for Two Years	<i>real speed, real trust</i>
04	Borrowed Is Shared	<i>the same well as every rival</i>
05	The Moat You Stopped Digging	<i>erosion is invisible until it isn’t</i>
06	Why Even ML Teams Get Seduced	<i>the trap is structural</i>
07	The Honest Counter-Case	<i>the quarter worth arguing</i>
08	The Verdict	<i>rent capability, build the moat</i>
09	Coda: Borrowing Without Dissolving	<i>the operating model</i>
—	Notes & Sources	<i>where the numbers come from</i>

Overture

The Two-Year Miracle

Read only this and you have the whole arc. Everything after it is evidence.

THE COMPANY

Start with a company that earned its place the hard way. It sold an AI product, and the reason customers paid for it was a decade of its own research: proprietary models trained on data nobody else had, a domain understanding that took years to accumulate, and an R&D group that was slow, expensive, and genuinely the source of the edge. The moat was not the brand or the sales team. The moat was the research, compounding quietly underneath the product.

THE REFRAME

Then foundation models arrived, and with them a product leader who preached letting go. Why fund a slow research lab when a public API is better than what we can train, ships next week, and costs a fraction of the salaries? Give up the heavy, expensive lab. Travel light on borrowed intelligence. Move at the speed of the market, not the speed of a PhD. It was not a stupid argument. For a while it was the right one.

THE MIRACLE

And it worked. Release cycles that used to take quarters took weeks. Costs fell. The roadmap filled with features the old lab would have called impossible on that timeline. Growth was real, customers were happy, and the CEO watched a leader turn a sleepy research-bound company into something that shipped. Trust compounded faster than the revenue did. The research group, now framed as legacy overhead, was quietly defunded. For two years, this looked like the best decision the company had made.

THE TURN

Then the third year arrived, and the ship started taking on water in a way nobody could quite explain, because nothing had broken. The product still worked. The model was still good. What changed is that everyone else had the same model. A competitor who started a year behind shipped the same feature in a weekend, because the intelligence was a public API for them too. The thing that used to take a decade of research to copy now took an afternoon of prompt engineering. The edge did not collapse. It evaporated, because it had been rented all along.

THE DIAGNOSIS

Here is the part worth sitting with. The product leader was not wrong that R&D was slow and expensive. They were wrong about what it was *for*. The research was not a cost center that happened to produce features. It was the moat-making machine, the one thing in the building that turned proprietary data and domain knowledge into an edge a rival could not buy off a price list. Borrowing intelligence got the company two years of speed by spending down an asset it had stopped replenishing. BCG's own breakdown of where AI value comes from puts roughly **70% in people, process, and proprietary data and only about 10% in the model itself.**¹ The leader optimized the 10% and dismantled the 70%.

The product did not fail because the AI was bad. It failed because the AI was borrowed. A leader traded the company's proprietary R&D for foundation-model speed, scaled brilliantly on capability every rival could rent, and called the missing moat "efficiency." You can rent capability. You cannot rent a moat.

THE THESIS, ~75% RIGHT, AND THE 25% WORTH ARGUING

The instinct to use borrowed intelligence was correct. Foundation models are extraordinary, and refusing them would have been its own failure. The error was treating a starting line as a finish line: borrowing the capability and never reinvesting a cent of the savings into the proprietary edge that borrowing was supposed to fund. Read on for what the moat actually was, and how to use the well without poisoning your own spring.

The Company That Had a Moat

Before the reframe, this company had something most AI startups never get: a reason it could not be cloned over a weekend.

It is worth being concrete about what the moat was, because the word gets thrown around loosely. This company's edge was not a single model. It was a loop. Years of customer usage produced proprietary data nobody else could legally or practically assemble. That data trained models tuned to a specific domain, which made the product measurably better at the one job customers cared about, which won more customers, which produced more proprietary data. The research group sat at the center of that loop, turning raw data into models and models into an edge.




Figure 1. The R&D flywheel. Proprietary data trains better models, which make a better product, which earns more usage and more proprietary data. The loop, not any one model, is the moat.

A flywheel like this is slow, and that slowness was the point. It took years to spin up, which is exactly why it took years for a competitor to match. A rival could copy the product's features in a quarter, but it could not copy the data flywheel underneath without running its own loop for the same years. That delay, measured in lead time over the competition, was the asset. Every turn of the wheel widened it.

I want to flag the trap before it springs, because it is the whole story. A flywheel that compounds invisibly also *erodes* invisibly. When it is spinning, you cannot see the work it is doing, so it is easy to mistake for overhead. That misreading is what the next section is about.

The Reframe

The product leader’s argument was not reckless. It was the most reasonable-sounding case in the room, which is precisely why it won.

Put yourself in that leadership meeting. A public foundation model has just beaten your in-house model on a benchmark your own researchers used to be proud of. It costs cents per call. It ships today. And here is your research group, asking for another year and another budget to maybe close a gap that an API just closed for free. The product leader stands up and says the quiet thing out loud: our R&D is no longer our edge, it is our anchor. Build on the frontier instead of trying to be it.

Every clause of that is locally true. The model really was better. It really was cheaper. The lab really was slow. What the reframe smuggled in was a hidden equivalence: that the model the lab produced and the moat the lab produced were the same thing. They were not. The model was the output. The moat was the loop that produced it, and the proprietary data and domain knowledge that fed the loop. You can replace the output with a better one off the shelf. You cannot buy the loop, which means when you shut the loop down to save money, you do not lose a model. You lose the ability to ever have a private one again.

WHAT THE REFRAME SAW	WHAT IT MISSED
R&D is slow and expensive.	Slow is what made it hard to copy.
The public model beats ours.	On the public benchmark. Not on your proprietary data.
We can ship features faster without it.	So can every rival, using the same API.
R&D is a cost center.	It was the moat-making machine.

This is the seduction, and it is worth naming as a pattern rather than a blunder. The reframe optimizes a real and visible cost while quietly retiring an asset whose value is real and invisible. On a quarterly dashboard, the cost shows up immediately and the lost moat shows up in three years. A leader rewarded on the dashboard will make this trade every time, and look brilliant for as long as the lag lasts.

Why It Worked for Two Years

The uncomfortable truth is that the strategy was not a slow-motion mistake. For two years it was a genuine, visible success.

I have to be fair to the leader here, because the story is useless if I pretend the borrowed approach never delivered. It delivered enormously. Freed from the cadence of a research lab, the product team shipped on the market's clock. Features that would have been a year of model work became a prompt and an integration. The cost line dropped as expensive researchers gave way to a usage-based API bill. Customers saw a product that suddenly moved fast, and they rewarded it.

And the most powerful output of those two years was not revenue. It was trust. The CEO had bet on a leader who turned a research-bound company into a shipping machine, and the bet was paying off in public. Every quarter of growth deepened the conviction that cutting R&D had been visionary rather than reckless. That compounding trust is the dangerous part, because it is what funds the next, larger cut. Success on borrowed capability does not just feel good. It buys the political capital to borrow harder.

THE LAG IS THE TRAP

Borrowed intelligence pays out immediately and bills you later. The speed is real in year one. The missing moat is also real in year one. You just cannot see it yet, and the visible win buys permission to ignore the invisible loss.

So when the turn came, it did not feel like a consequence of the original decision, because two years and a wall of good results sat between cause and effect. That distance is exactly why this failure mode is so hard to catch in the act. The reward and the bill arrive on completely different schedules.

Borrowed Is Shared

Here is the structural fact the speed obscured: capability you rent is capability your competitor can rent on identical terms.

A foundation model is a shared well. The whole point of it, the reason it is so cheap and so good, is that its cost is spread across every company that calls it. That is wonderful when you are deciding whether to use it. It is fatal when it is the *only* intelligence in your product, because the input that powers your differentiation is available, at the same price and quality, to everyone you compete with. You cannot build a height advantage on a floor that is shared by the whole market.




Figure 2. Borrowed intelligence is shared intelligence. Every rival builds on the same slab, so the input that powers your product cannot, by itself, make it taller than theirs for long.

This is why the differentiation question has gotten so sharp. When a single API delivers most of the capability, the model is no longer where you compete; it is table stakes. The defensibility has to come from somewhere the well cannot reach: proprietary data the model has never seen, a workflow tuned to a domain over years, switching costs, distribution, trust. That is the same conclusion BCG reaches from the cost side, where roughly 70% of realized AI value sits in proprietary data and process rather than the model.¹ The capital markets noticed the gap too. Sequoia’s widely cited “\$600B question” pointed out how much AI revenue would have to materialize to justify the infrastructure spend, and much of that revenue was being booked by companies whose only asset was a thin layer over someone else’s model.²

A rented input cannot be a private advantage. The well is shared, and so is

whatever you draw from it.

The Moat You Stopped Digging

A moat does not collapse the day you stop maintaining it. It fills in slowly, invisibly, until the morning a competitor walks straight across.

The cruelest property of this failure is its timing. When the company stopped running its flywheel, its lead over competitors did not vanish. It started decaying, one turn of the wheel at a time, while the dashboard showed nothing but green. The proprietary data stopped accumulating an edge. The domain models stopped improving past the public frontier. Rivals, running their own borrowed-intelligence playbook, kept closing the distance the old R&D used to keep open.

YOUR LEAD OVER RIVALS, WITH THE FLYWHEEL SWITCHED OFF




Figure 3. Lead time is a moat only while you keep extending it. Borrowed capability does not extend the lead; rivals calling the same model close it on their own schedule.

By the time the damage shows up in the numbers, the recovery is brutal. Restarting a research flywheel is not a budget you switch back on. The team has scattered, the tooling has rotted, and the data advantage you would train on is the very thing you let go stale for three years. The company that took a decade to build its moat now has to rebuild it from behind, against rivals who never had to, while explaining to a board why the strategy that looked brilliant for two years suddenly is not. The asset was cheap to spend and expensive past the point of pain to replace.

This is the real meaning of the sinking ship. Nothing dramatic breaks. The product keeps working, the model keeps responding, the demos keep demoing. What is gone is the reason anyone had to choose you over the rival with the same API, and that reason was the quiet flywheel the company decided was overhead.

Why Even ML Teams Get Seduced

If this only happened to companies that never understood AI, it would be a competence story. It happens to the teams who understand it best, which makes it a structural one.

The temptation is to read this as a product leader overruling the smart researchers. The harder truth is that the researchers often agree, at least at first, and for good reasons. A strong ML team is the first to admit when a public model genuinely beats their in-house one, because they respect the benchmark. They are tired of defending a slow roadmap. And borrowed capability lets them ship things they are proud of without a year of grinding. The seduction works on the experts precisely because they are honest about what the frontier models can do.

The structural forces line up the same way every time, which is why I treat this as a pattern rather than a lapse of judgment:

PULLING TOWARD BORROWING

- The cost saving is immediate and shows on the dashboard.
- The speed win is visible to customers this quarter.
- The public model really is better on public benchmarks.
- R&D's value does not show up on a quarterly review.

PULLING TOWARD THE MOAT

- The lost edge is real but invisible for years.
- Proprietary data compounds only if you keep feeding it.
- The benchmark is public; your data is not.
- Defensibility shows up in churn and pricing power, late.

Read the two columns and the outcome is almost mechanical. Every force on the borrowing side is fast, visible, and rewarded. Every force on the moat side is slow, invisible, and easy to defer. An organization that runs on quarterly evidence will tilt toward borrowing even when its smartest people know better, because the scoreboard only measures one column. This is why “just hire great ML people” does not protect you. The trap is not ignorance. It is an incentive structure that pays out for spending the moat and stays silent while it erodes.

The Honest Counter-Case

A thesis you cannot argue against is not strong, it is unexamined. Here is the quarter of this argument I hold loosely.

Not all R&D is a moat

Plenty of corporate research is vanity, defended by sunk cost and prestige. Some labs produce papers and not edges. If a company's R&D was chasing a frontier the public models now own outright, cutting it was not the crime; clinging to it would have been. The argument is not "never cut research." It is "know whether the research you are cutting is the moat or the vanity," and that is a hard, specific judgment, not a slogan.

Borrowing genuinely raises the floor

Foundation models democratized capability that used to be locked behind nine-figure research budgets. A small team can now do things that were impossible a few years ago. Refusing to use borrowed intelligence on principle would be its own failure mode, slower and prouder and just as fatal. The point is not to avoid the well. It is to not mistake the well for a private spring.

The moat can legitimately live elsewhere

Model R&D is one moat, not the only one. A company can defend itself with proprietary data, distribution, brand, regulatory position, switching costs, or a workflow lock-in that no API touches. Some companies are right to rent the intelligence and build their moat entirely in distribution or data. The failure in this story was not borrowing the model. It was borrowing the model and building no moat *anywhere* with the time and money that borrowing freed up.

"Borrow then build" is a valid sequence

The strongest version of the product leader's plan is one I would have backed: borrow to move fast now, and reinvest the savings into a proprietary edge that compounds later. Borrowing to bootstrap is good strategy. The error was stopping at the first half, treating a starting line as a finish, and pocketing the savings as margin instead of plowing them into the moat the borrowing was meant to fund.

None of this softens the core claim. It sharpens it. Borrowed intelligence is a fine input and a fatal sole asset. The two-year miracle was real, and so was the bill. Hold the thesis at 75%, allow that some research

deserved to die and some moats live outside the lab, and the argument gets stronger: the failure was never using the well, it was forgetting to dig.

Rent the Capability, Build the Moat

So where does this leave a leader staring at a brilliant, cheap, public model and a slow, expensive lab?

Not at “refuse the model,” and not at “kill the lab.” The model is a gift; take it. The mistake was never the borrowing. It was treating borrowed capability as the whole strategy instead of the floor you build on. Stated as plainly as I can:

THE VERDICT

Borrowed intelligence is a *starting line*, not a strategy. Rent the commodity capability that every rival can also rent, and spend what you save building the one thing they cannot: a proprietary edge that compounds, fed by data and domain work only you can do. The leaders who sink treat the rented floor as the finished building. The ones who last treat it as the floor, and build the moat on top.

Notice what that does not say. It does not say train your own foundation model; almost no one should. It does not say research is sacred; some of it deserved to die. It says the savings from borrowing are not margin to harvest, they are capital to reinvest, and the place to reinvest is the proprietary layer the well cannot reach. The reason the company sank was not that it borrowed. It is that it borrowed and never built.

WHERE THE MOAT ACTUALLY LIVES




Figure 4. Rent the commodity floor. Spend what you save building the layer competitors cannot rent. The moat was never the model.

The instinct to borrow was right. The correction is the whole point: a rented input is a starting line, the model was the easy part, and the moat is the proprietary loop you build on top of it. Use the well. Keep digging your own.

Borrowing Without Dissolving

It is one thing to say rent the floor and build the moat. It is another to run it on Monday without lying to yourself about which is which.

Six rules for using borrowed intelligence

01 Separate the commodity layer from the moat layer, explicitly.

Draw the line on a whiteboard: what you rent because everyone can, and what you build because only you can. If you cannot name the moat layer, you do not have one yet.

02 Treat borrowing's savings as reinvestment capital, not margin.

The money and time you save by renting the model are the budget for the proprietary edge. If the savings go straight to the bottom line, you are harvesting the moat, not building it.

03 Keep the data flywheel turning even when the model is rented.

You can borrow the intelligence and still own the proprietary data, the labels, the feedback, and the domain evaluations the public model never sees. That ownership is the moat that survives the next model release.

04 Measure your lead, not just your shipping speed.

Track time-to-copy: how long would it take a well-funded rival with the same API to match this feature. If the answer is a weekend, you shipped a commodity, however fast you shipped it.

05 Fund R&D as moat maintenance, with a moat metric.

Do not defend research on prestige. Defend the specific work that widens an edge a rival cannot rent, and kill the work that does not. Tie the budget to defensibility, not to publications.

06 Watch the trust the wins are buying.

A run of borrowed-capability success buys the political capital to borrow harder. Name that out loud at the top, so a two-year winning streak does not quietly fund the cut that sinks

you.

Frontier and bedrock

One way to keep the plan honest is to separate where you actually compete from what is simply true regardless of who wins.

FRONTIER • WHERE YOU EARN AN EDGE

- Proprietary data the public model has never seen and cannot get.
- Domain evaluations and feedback loops only your usage produces.
- Workflow and integration lock-in that no API replicates.
- A flywheel that compounds an advantage rivals must spend years to match.

BEDROCK • TRUE REGARDLESS OF WHO WINS

- The frontier model is rentable by every competitor at the same price.
- Borrowed capability scales fast and differentiates not at all.
- Moats erode invisibly and rebuild expensively.
- The model is the cheap 10%; the moat is the 70% you have to own.

Read the two columns together and the strategy falls out on its own. You spend your scarce capital on the frontier column, and you stop pretending the bedrock column is negotiable, because those are facts you build on rather than bet against. The leaders who survive the next three years will be the ones who borrowed the floor without mistaking it for the building, and kept digging the moat while the borrowing looked, for a while, like it had made the moat unnecessary.

*Rent the intelligence. Own the reason
anyone needs you to.*

Appendix

Notes & Sources

The company arc in this report is a composite, drawn from the recurring pattern I have watched across AI product organizations, not a single named firm. The analytical claims are grounded in the public sources below; citation numbers in the text map to them.

Figures are headline findings from the named institutions; exact wording varies by source release.

1. Boston Consulting Group – *Where's the Value in AI?* (2024). The 10-20-70 split of realized AI value: roughly 10% algorithms, 20% technology and data, 70% people and process. Finding: a minority of companies move beyond proofs of concept to durable value.
2. Sequoia Capital (David Cahn) – *AI's \$600B Question* (2024). On the gap between AI infrastructure spend and the revenue required to justify it, and the weak defensibility of thin application layers built over shared foundation models.
3. MIT (Project NANDA) – *The GenAI Divide: State of AI in Business 2025*. Finding: ~95% of organizations report no measurable P&L return on generative-AI pilots, consistent with borrowed-capability bets that fail to compound into an edge.
4. General industry analysis on AI defensibility and “thin wrapper” economics (a16z, Sequoia, and others, 2023-2025) – on proprietary data, distribution, and workflow as the durable moats once the model itself is a shared commodity input.
5. Public model-capability tracking (open-weight models narrowing the gap to closed frontier models, 2024-2025) – context for why a rented capability edge erodes as the frontier becomes broadly and cheaply available.

© 2026 Anuj Sadani. A field report on borrowed intelligence and the moats it dissolves. Read the Overture for the arc; read the sections for the evidence.