

Three Leaders, One *Market*

Three experienced leaders, three good-faith strategies, all backed by data and the board. The AI era broke all three, by different paths. There are no villains here, only the people left on the stake.

TO **Boards, CEOs, and the leaders carrying the AI bet**
FROM **Anuj Sadani** · Principal SDE, builder of AI systems and the teams behind them
RE **Why three sincere AI strategies all failed, and who pays**

THE STEWARD

“The proven playbook.”

held the line too long

THE SURFER

“Ride every wave.”

nothing compounded

THE ARCHITECT

“Build the right system.”

the ground moved

At a glance. Three sincere leaders, three good-faith strategies, one market that kept changing the rules. Clay is the Steward, gold the Surfer, teal the Architect. All three failed.

Read the **Overture** in five minutes for the whole story. Read the rest for each leader’s honest case, the one thing that moved under all of them, and why the failure was not a character flaw.

By **Anuj Sadani** · asadani.tech · June 2026

CONTENTS

—	Overture: Three Good Bets, One Moving Target	<i>the story in five minutes</i>
01	The Same Starting Line	<i>all three were sound</i>
02	The Steward, Who Held the Line	<i>the proven playbook</i>
03	The Surfer, Who Rode Every Wave	<i>change without compounding</i>
04	The Architect, Who Built the Right System	<i>overtaken, not wrong</i>
05	The Moving Target	<i>what changed under all of them</i>
06	Three Failure Paths, One Cause	<i>a fixed stance on change</i>
07	The Honest Counter-Case	<i>maybe one was unlucky</i>
08	The Verdict	<i>no villains, still on the stake</i>
09	Coda: On the Stake	<i>holding a moving target</i>
—	Notes & Sources	<i>where the numbers come from</i>

Overture

Three Good Bets, One Moving Target

Read only this and you have the whole story. Everything after it is evidence, and none of it is an accusation.

THREE SINCERE BETS

Three companies in the same market, the same years, the same pressure to do something real about AI. Three leaders, each experienced, each backed by data and a board that signed off. None of them was careless. Each chose a strategy that had a strong case at the time, and each gave it their conviction. I want to say that plainly at the start, because the easy version of this story looks for the fool, and there is no fool in it.

THE STEWARD

The Steward ran the proven playbook: the structured AI program, the center of excellence, the governance, the disciplined pilots. This was not recklessness. It was the widely endorsed best practice, the same approach that had carried the company through every prior technology wave, blessed by consultants and benchmarks. The Steward did not change it, because for twenty years changing it had been the mistake.

THE SURFER

The Surfer did the opposite, and could defend every move. Each quarter brought a better model, a new paradigm, a fresh benchmark, and the Surfer re-pointed the company at it. The board loved the energy and the press releases. Staying current looked like exactly the right instinct in a field moving this fast. The Surfer changed constantly, because standing still looked like the real risk.

THE ARCHITECT

The Architect made one deep, careful bet: rebuild the operating model around a specific, well-reasoned thesis about where value would sit, and let it compound. This was the most sophisticated of the three, the one most people would call correct. The Architect committed once and built for the long run, because that is how durable advantage has always been made.

THE TURN

All three failed. Not equally, and not on the same schedule, but all three. The Steward's playbook was built for a technology whose rules stayed put, and these rules did not. The Surfer's constant motion meant nothing ever compounded into an advantage. The Architect committed correctly to a target that then moved, and watched a carefully built lead get overtaken before it paid. The common thread is not incompetence. It is that the ground itself kept shifting, faster than any one fixed strategy could track.

Three experienced leaders, three strategies that had each worked somewhere before, each backed by data and signed off by a board. One held the proven line, one rode every wave, one built the careful system. The era broke all three, because the ground they stood on kept moving. There are no villains here. There is only the fact that when the company fails, the leader is the one left on the stake.

THE THESIS, ~75% RIGHT, AND THE 25% WORTH ARGUING

This is not a piece about who was stupid. It is a piece about a market that kept changing what it rewarded, and about three good people who each held one honest, defensible position against a target that would not hold still. Read on for each of their cases, told fairly.

§ 01 • The Control

The Same Starting Line

Before the three diverged, it is worth establishing that none of them started from weakness or bad faith.

Each leader had real credentials and a real mandate. The same access to frontier models on the same terms. Comparable budgets, capable teams, and boards that had approved the plan after seeing the data behind it. Each strategy came with evidence and a respectable precedent. If you had sat in any of the three board meetings, you would have nodded along, because each case was sound by the standards of its moment.

I labor this because the reflex when a company fails is to find the flaw in the person who ran it. Sometimes that flaw is there. Often it is not, and the search for it is a way of avoiding a harder truth: that a competent, well-resourced, well-intentioned leader can execute a defensible strategy and still lose, because the conditions the strategy assumed quietly stopped being true. That is the situation all three were in, and pretending otherwise would be the dishonest part.

THE CONTROL GROUP

Same market, same models, same budgets, same board approval, three sincere strategies. When three good-faith bets all fail, the cause is unlikely to be three independent character flaws. Look at what they were all standing on instead.

§ 02 • The Steward

The Steward, Who Held the Line

The Steward ran the strategy that every responsible playbook recommended. That is exactly why it is the hardest failure to see coming.

Give the Steward full credit, because the approach was textbook and the textbook was good. A clear AI strategy, a center of excellence to concentrate scarce talent, governance so the company would not do something reckless, and a disciplined sequence of pilots gated on return before anything scaled. This is the method that absorbed cloud, mobile, and enterprise software into the business without disaster. Consultants endorsed it. Peer companies ran the same play. The board approved it precisely because it was the safe, proven path.

It failed for a reason that had nothing to do with effort. The proven playbook assumes the technology behaves like the last one: fixed at purchase, measurable in a quarter, scaled once the return is clear. A capability that improves only with use, whose value the market keeps redefining, does not fit that frame, so the gates rejected the very bets that needed time, and the program produced motion rather than shipped decisions. This is the base rate, not bad luck. MIT's 2025 study found **95% of organizations getting no measurable return** on their generative-AI pilots.¹ The Steward was median, running the method that used to guarantee not being median.

The cruelty of it is that the Steward could not have been faulted in advance. Holding the proven line is what good stewardship had always meant. The line simply stopped being where the value was.

§ 03 • The Surfer

The Surfer, Who Rode Every Wave

The Surfer did the thing everyone says to do in a fast-moving field: stay current. It still did not work.

The Surfer's instinct was defensible and, in isolation, admirable. AI was moving faster than any technology in memory, so the Surfer kept the company on the frontier: adopting each new model the week it shipped, re-architecting around the latest paradigm, chasing every benchmark that moved. There was always fresh data to justify the next pivot, and a

board that read the same headlines and applauded the company for not falling behind. Energy was high. The roadmap was always exciting.

The failure was not in any single move. Each was reasonable. The failure was that constant change let nothing compound. A data flywheel needs time turning to build an edge. An evaluation suite gets valuable only as it accumulates. A workflow earns trust by staying stable long enough for people to rely on it. The Surfer reset all of that every quarter, so the company was permanently three months into everything and finished with nothing. The team developed whiplash, the budget drained into migrations, and the one thing the constant motion never produced was a durable advantage. Worth remembering that BCG put roughly **70% of AI value in people, process, and proprietary data**, the parts that only compound with time, and only about 10% in the model itself.² The Surfer kept swapping the 10% and never let the 70% accrue.

Calling the Surfer a hype-chaser is the easy verdict and the wrong one. The Surfer was responding rationally to a real signal, that the frontier was moving fast. The mistake was treating speed of change as the thing to match, rather than something to absorb selectively. Sincere, defensible, and fatal.

§ 04 • The Architect

The Architect, Who Built the Right System

The Architect made the bet I would probably have made. That is what makes this one sting.

The Architect did the sophisticated, patient thing. Rather than announce or chase, the Architect formed a clear thesis about where durable value would sit, then rebuilt the operating model around it: redesigned workflows, an evaluation culture, proprietary data flywheels, accountability written down before launch. This is the approach I have argued for, the one that compounds while rivals run pilots. For two years it looked like the obvious winner, and by the internal evidence it was working.

It failed anyway, and not because the execution was poor. The Architect committed to a specific picture of the future, and the future moved. A model release commoditized part of what the company had carefully built. The market quietly re-decided what counted as defensible, and the thesis the whole system was optimized around was no longer the thesis that mattered. The Architect had built the right system for a target that shifted while it was being built. Deep commitment is how durable advantage is made, and it is also what makes you brittle when the thing you committed to is overtaken.

RIGHT EFFORT, MOVED TARGET

The Architect's failure is the most uncomfortable, because nothing was done wrong by the understanding of the day. A correct, well-built bet can still be overtaken when the market changes what it rewards. Conviction and brittleness are the same trait seen from two sides.

Of the three, the Architect is the one most people would have backed, which is the point. This was not the failure of a weak strategy. It was the failure of a strong one aimed at a target that did not stay put.

The Moving Target

Hold the three failures together and the shared cause is not in any of the leaders. It is in the ground they stood on.

What actually changed underneath all three was the market's definition of winning. The bar did not just rise, it kept relocating. Early on, having a chatbot at all was credible. Then the credible thing was retrieval and copilots. Then it was agents in production. Then it was a defensible edge a foundation model could not erase. Each shift invalidated a strategy that had been correct against the previous bar, and each shift arrived before the last one had been fully answered.




Figure 1. The market kept moving what it rewarded. A strategy aimed at the 2023 bar was pointed at a spot the target had already left by the time it shipped.

This relocation was not gentle, and it was widely felt. The share of enterprises chasing agentic approaches jumped from roughly a quarter to two-thirds in a single year as the consensus moved.⁴ By 2025, **42% of companies had abandoned most of their AI initiatives**, up from 17% a year earlier, much of it the sound of strategies being scrapped because the target had moved out from under them.³ When the definition of success changes every year, no fixed strategy is safe, because the thing that made it correct keeps expiring.

The goalposts did not only move. They moved faster than any one bet could be carried to them.

Three Failure Paths, One Cause

Line the three up and the structure is clear: each chose a fixed relationship to

change, and the target was not fixed.

The Steward chose almost no change, trusting a proven method. The Surfer chose constant change, trusting motion. The Architect chose a single deep change, trusting one well-reasoned commitment. Three different stances, and the market punished all of them, because it did not reward a stance toward change at all. It rewarded continuous re-grounding, which is not a fixed position any of them could have picked once and held.




Figure 2. Three fixed stances on one axis. The cadence the era rewarded was not a point on it, and it would not hold still, so no single choice could track it.

	THE STEWARD	THE SURFER	THE ARCHITECT
Stance on change	held the proven line	changed every quarter	one deep commitment
It was backed by	consultants, precedent	the latest benchmark	a reasoned thesis
How it failed	the playbook expired	nothing compounded	the target moved
The flaw, fairly	none, until the rules changed	none, except no patience	none, except brittleness

Read across the bottom row, because it is the honest one. The flaw in each case is “none,” with a qualifier that only becomes visible after the ground moves. That is what makes this an era problem and not a competence problem. Each leader optimized correctly for a world that did not stay long enough to reward them.

§ 07 • The Counter-Case

The Honest Counter-Case

A story where everyone is blameless and the universe is at fault should make you suspicious. Here is the quarter of it I hold loosely.

Sometimes it really is the leader

I have leaned hard on “no villains,” and I should be honest that some failures are not like this. There are leaders who do not understand the technology, who chase prestige, who will not listen to their own engineers. Those failures are about the person. My claim is only that a great many AI failures are not those, and that the reflex to assume they are does real

harm to good people.

One of them might have made it

A story where all three fail is tidy, and reality is lumpier. In a different draw, the Architect's bet pays before the ground moves and looks like genius. The Steward's discipline is exactly right in a market that happens to settle down. Outcomes carry more luck than we like to admit, and I do not want the clean structure to hide how much of this is timing rather than logic.

“Re-ground continuously” is easy to say

My implied prescription, track the moving target, is close to unfalsifiable as advice. Continuous re-grounding can curdle into the Surfer's thrash just as easily as it can save you, and from inside the moment, the moving target and the passing fad look identical. I am describing the problem more confidently than anyone can prescribe the solution, and pretending otherwise would be its own dishonesty.

Some of this is survivorship

The companies that won this market did some of the same things these three did, and we credit their judgment after the fact. The line between the Architect who was overtaken and the architect who was vindicated is often drawn by events neither could control. Be careful taking too clean a lesson from three losses without weighing the wins that looked identical on the way up.

None of this rescues any single strategy. It softens the verdict on the people. The core holds: the market kept moving what it rewarded, three sincere bets were aimed at a target that would not hold still, and the failures were structural more than personal. Hold it at 75%, allow for genuine fault and genuine luck, and the humane reading survives.

§ 08 • The Verdict

No Villains, Still on the Stake

So name it plainly, with the dignity these three are owed and the honesty the outcome demands.

None of the three was foolish or lazy. Each brought experience, conviction, and a defensible read of the data, and each had the approval to proceed. The market simply kept changing what it rewarded, faster than any fixed strategy could be carried to it. That is the real condition of this era, and it is not a flattering one for the idea that the right plan, well executed, is enough.

*Three good-faith strategies, three honest failures, one market that would not hold still. The Steward held a line that expired, the Surfer chased a target and never compounded, the Architect built the right system for a future that moved. There are **no villains** in this story. There is only the fact that when the company fails, the leader is the one left on the stake.*




Figure 3. Three paths, one outcome. The Steward holds flat, the Surfer spikes and decays, the Architect rises then is overtaken. The bar the market rewarded kept climbing away from all three.

The point of the stake is not blame. It is responsibility, which is different and heavier. A leader signs for the bet, carries the team's belief, and absorbs the cost when the conditions betray the plan. That asymmetry is the job. It is worth naming out loud, because the people who do this work deserve to have the difficulty of it stated honestly, rather than flattened into a hunt for what they did wrong.

§ 09 • Coda

On the Stake

If you are the one carrying the AI bet right now, the useful question is not which of the three you resemble. It is how to hold a target that will not hold still.

I cannot offer a strategy that beats a moving market, because a fixed one cannot, and that is the whole lesson. What I can offer is a way of standing that is less brittle than any of the three, drawn from watching all three break.

01 Separate the bet from the bar.

Write down, explicitly, what the market currently rewards and re-check it on a schedule. Most strategies fail not because the bet was wrong but because nobody noticed the bar had moved.

02 Let some things compound while you re-ground others.

The Surfer's lesson is that not everything can be in motion. Keep the data, evaluations, and judgment compounding, and re-aim the strategy on top of them. Stability underneath, adaptation on top.

03 Hold convictions loosely enough to be overtaken gracefully.

The Architect's lesson is that deep commitment and brittleness are the same trait. Build so that being wrong about the thesis costs you a pivot, not the company.

04 Name who is on the stake, and protect them.

The leader carrying the bet absorbs the cost when conditions change. Boards that understand this judge the decision in its context, not the outcome in hindsight, and they keep good people through an honest loss.

None of that guarantees survival, and I will not pretend it does. The Steward, the Surfer, and the Architect were not lesser versions of some wiser leader who would have gotten it right. They were three good people who each held one honest position against a market that kept changing the rules. When you read about a company's AI strategy failing, look first at what moved under it before you look for the fool. Usually the ground moved, and someone competent was standing on it.

*There are no villains in a market that keeps
changing the rules. There is only whoever
signed the strategy when the ground gave
way.*

Appendix

Notes & Sources

The three leaders are composites, drawn from recurring patterns I have watched across AI programs, not three named firms. The cited statistics are public findings; citation numbers in the text map to them, and exact wording varies by source release.

1. MIT (Project NANDA) – *The GenAI Divide: State of AI in Business 2025* (August 2025): ~95% of organizations report no measurable P&L return on generative-AI pilots. nanda.media.mit.edu
2. Boston Consulting Group – *Where's the Value in AI?* (2024): the 10-20-70 split of realized AI value (10% algorithms, 20% technology and data, 70% people and process). bcg.com

3. S&P Global Market Intelligence – enterprise AI survey (2025): 42% of companies abandoned most of their AI initiatives, up from 17% the prior year. spglobal.com
 4. Analyst surveys on agentic-AI adoption (2023–2025) and the Gartner AI Hype Cycle – the year-over-year shift in what enterprises pursued (chatbots, then retrieval, then agents, then defensibility), evidence that the market’s definition of success kept relocating. gartner.com
-

© 2026 Anuj Sadani. A story on three good-faith AI strategies that all failed, and the moving target underneath them. Read the Overture for the argument; read the sections for the evidence.